

1 Use of general purpose graphical processing units with 

2 MODFLOW-2005

3 Hughes, J.D.

White, J.T.

U.S. Geological Survey

U.S. Geological Survey

Tampa, Florida

Tampa, Florida

`jd Hughes@usgs.gov`

`jwhite@usgs.gov`

4 May 18, 2012

5 **Abstract**

6 To evaluate the use of general-purpose graphics processing units (GPGPU) to im-
7 prove the performance of MODFLOW-2005, an unstructured preconditioned conjugate
8 gradient (UPCG) solver has been developed. The UPCG solver uses a compressed
9 sparse row storage scheme and includes Jacobi, zero fill-in incomplete and modified-
10 incomplete LU factorization, and generalized least-squares polynomial precondition-
11 ers. The solver utilizes the NVIDIA CUDA-enabled implementation of the Basic Lin-
12 ear Algebra Subprograms (BLAS) library for calculations performed on the GPCPU
13 (CUBLAS). The UPCG solver also includes options for sequential and parallel solu-
14 tion on the central processing unit (CPU) using OpenMP. For simulations utilizing the
15 GPGPU, all dot-product, element-by-element vector operations, and matrix-vector op-
16 erations are performed on the GPGPU; memory copies between the central processing
17 unit CPU and GPCPU occur prior to the first iteration of the UPCG solver and after
18 satisfying user-specified infinity-norms for head and flow or a user-specified maximum

number of iterations is exceeded. Because of the sequential nature of application of the incomplete LU factorization preconditioners, these calculations are performed on the CPU.

The efficiency of the UPGP solver for GPGPU and CPU solutions is benchmarked using simulations of a synthetic, heterogeneous unconfined aquifer with tens of thousands to millions of active grid cells. Testing indicates GPGPU speedups on the order of 2-8, relative to the standard MODFLOW-2005 PCG solver, can be achieved when (1) memory copies between the CPU and GPGPU are optimized, (2) the percentage of time performing memory copies between the CPU and GPGPU is small relative to the calculation time, (3) high-performance GPGPU cards (*e.g.*, NVIDIA FERMI architecture) capable of high-speed memory access and double-precision calculations are utilized, and (4) CPU-GPGPU combinations are used to execute sequential operations that are difficult to parallelize. Furthermore, UPGP solver testing indicates GPGPU speedups exceed parallel CPU speedups achieved using OpenMP on multi-core CPUs for preconditioners that can be parallelized.

1 Introduction

MODFLOW-2005 (Harbaugh 2005) is a finite-difference groundwater flow model that has been effectively applied to two- and three-dimensional problems since 1984. Because MODFLOW-2005 uses a cell-centered finite-difference (CCFD) approximation and a rectilinear grid, model runtimes are often increased by factors of two, four, and eight with increased vertical, horizontal, or combined vertical and horizontal discretization, respectively. To date, larger model sizes have been accommodated through a combination of faster CPUs and better linear solvers. The original version of MODFLOW (McDonald and Harbaugh 1988) included 1) a strongly implicit procedure (SIP) linear solver and 2) a slice-successive overrelaxation (SOR) linear solver. These original linear solvers were supplemented with the preconditioned conjugate gradient (PCG2) linear solver (Hill 1990), the link-algebraic multi-grid (LMG) linear

solver (Mehl and Hill 2001), the geometric multigrid (GMG) linear solver (Wilson and Naff 2004), and the preconditioned conjugate gradient with improved nonlinear control (PCGN) linear solver (Naff and Banta 2008).

Parallelization of linear solvers is another approach for accommodating larger model sizes. Dong and Li (2009) parallelized the MODFLOW PCG solver (Hill 1990) using the OpenMP (OpenMP Architecture Review Board, 2005) programming paradigm and observed speedups (CPU runtime / parallel runtime) as large as 5 on machines with multi-core CPUs. Naff (2008) developed a parallelized linear solver for MODFLOW using the Message Passing Interface (MPI) standard and observed a speedup of 7. In addition to OpenMP and MPI, GPGPUs also presents an option for linear solver parallelization. In comparison to CPUs, GPGPUs are much faster, have higher bandwidth, and typically more cores (*e.g.*, 448 in the case of the NVIDIA[®] Tesla[™] C2050 and C2070 GPGPUs). Use of GPGPUs to parallelize linear solvers for MODFLOW-2005 is also attractive because of the availability of the CUDA (Compute Unified Device Architecture) application programming interface (API) for NVIDIA GPGPUs, which supports development of parallel GPGPU code in C/C++ or Python, Fortran, Java, and MATLAB using wrappers around the C/C++ implementation. For example, Soni *et al.* (2012) used the CUDA language to develop a GPGPU code to solve computational fluid dynamics problems. For solution of the groundwater flow equation, Ji *et al.* (2012) developed a parallel GPGPU version of MODFLOW-2000 with a Jacobi preconditioner and observed speedups ranging from 1.8 to 3.7.

In this paper, we evaluate speedups resulting from GPGPU parallelization of the unstructured UPCG linear solver for a unconfined groundwater flow problem. We also evaluate the performance of the UPCG linear solver using Jacobi, zero fill-in modified incomplete LU (MILU0), and generalized least-squares polynomial (GLSPOLY) preconditioners on the CPU and GPGPU. And finally, we evaluate the performance of the UPCG solver on the GPGPU to parallel CPU simulations using the OpenMP parallel programming paradigm.

2 GPGPU Parallel Programming

GPGPU code developed using the NVIDIA[®] CUDA API is essentially a sequential code that is capable of using the fork-join model of parallel execution similar to OpenMP. In a fork-join model, a single master thread is active when program execution begins. The master thread executes sequential portions of the code. At points where parallel operations are required, the master thread spawns, or forks, additional threads that work concurrently with the master thread through the parallel section. At the end of the parallel code, the spawned threads are suspended and rejoined to the master thread. The fork-join model used in OpenMP is shown graphically in **figure 1a**. Typically, the number of OpenMP threads used in the fork-join model is less than the total number of cores available on multi-core CPUs.

A GPGPU is a set of streaming multiprocessors (SM) with each having many CUDA cores (CC) and each CC is capable of executing many threads simultaneously. Each SM includes on-chip register, shared, constant cache, and texture cache memory. The GPGPU also includes uncached device memory with higher latency and lower bandwidth than on-chip memory (NVIDIA, 2011).

For GPGPU code implementations, the main code is a sequential FORTRAN/C++ code which forks the data into many threads at the invocation of a CUDA kernel. The CUDA API is a software environment that allows C/C++ functions, called kernels, to be developed that run on a GPGPU. At runtime, thousands of GPGPU threads are generated by a CUDA kernel that are executed on a CUDA grid. The CUDA grid is composed of one or more blocks that are executed on CCs on a SM. There is a limit to the number of threads that can be run in a single block which depends on the GPGPU unit model. The number of blocks used for a given problem is calculated as,

$$\text{numBlock} = \frac{(\text{NCELL} + \text{threadsPerBlock} - 1)}{\text{threadsPerBlock}}, \quad (1)$$

where `NCELL` is the number of active model cells and `threadsPerBlock` is the number of threads that can be run in a single block. For increased throughput, threads within a block are further split into warps with a maximum number of threads that are managed and processed simultaneously on CCs. Multiple warps execute one after another to complete block processing. All CCs within a SM share a single instruction, which results in massively parallel processing potential. GPGPU architecture is shown graphically in **figure 1b**.

In a kernel, registers are used to hold frequently accessed variables that are private to each thread. Shared memory is allocated to blocks and threads in one block can only access shared memory of their own block. **Additionally, all threads can access global memory.**

A NVIDIA Tesla C2050, based on the NVIDIA Tesla T20 GPU, was used to evaluate the UPGP solver. The Tesla C2050 has the following specifications: SMs = 14, CCs per SM = 32, `threadsPerBlock` = 1,024, and warp size = 32. The multi-threaded instruction unit on each of the 14 Tesla C2050 SMs can run 48 warps concurrently, yielding a total of 21,504 concurrent threads (14 SMs \times 48 warps per SM \times 32 threads per warp = 21,504 concurrent threads). The total number of concurrent threads that can be run on a GPGPU greatly exceeds the number OpenMP threads that can be run in parallel on today's multi-core CPUs. For example, it is possible to purchase server based systems with 8 Intel E7-8800 CPUs having 10 cores per CPU but even this system would have maximum thread counts more than two orders of magnitude less than available with the NVIDIA C2050 GPGPU. With thread counts on the order of tens of thousands, GPGPUs have great potential for parallelization of MODFLOW-2005.

3 Conjugate Gradient Linear Solver

The constant-density, three-dimensional groundwater flow equation is described by the partial-differential equation,

$$\frac{\partial}{\partial x} \left(K_{xx} \frac{\partial h}{\partial x} \right) + \frac{\partial}{\partial y} \left(K_{yy} \frac{\partial h}{\partial y} \right) + \frac{\partial}{\partial z} \left(K_{zz} \frac{\partial h}{\partial z} \right) + W = S_S \frac{\partial h}{\partial t}, \quad (2)$$

119 where K_{xx} , K_{yy} , K_{zz} are the hydraulic conductivity [L/T] along the x, y, and z-coordinate
 120 axes which are assumed to be parallel to the major axes of hydraulic conductivity; h is
 121 the groundwater head [L]; W is a volumetric flux per unit volume [L³/L³T] representing
 122 sources/sinks of water; S_S is the specific storage [L⁻¹]; and t is time [T].

123 MODFLOW-2005 uses a cell-centered, finite-difference approximation of **equation 2** to
 124 solve for three-dimensional groundwater flow (Harbaugh 2005). Development of the finite-
 125 difference form of **equation 2** using the continuity equation for a cell results in

$$\sum Q = S_S \frac{\Delta h}{\Delta t} V, \quad (3)$$

126 where Q is the flow rate [L³/T] into a cell from adjacent cells and sink/source terms, Δh
 127 is the head change [L] in a cell over a time interval Δt [T], and V is the volume of a cell
 128 [L³]. Manipulation of **equation 3** to separate known and unknown h terms results in a large
 129 system of simultaneous linear equations of the form,

$$\mathbf{A}\mathbf{h} = \mathbf{b}, \quad (4)$$

130 where \mathbf{A} is a square, symmetric, positive-definite coefficient matrix [L²/T] that includes cell-
 131 by-cell conductances, calculated using cell dimensions and hydraulic conductivities at cell
 132 interfaces, and unknown components of sink/source and storage terms; \mathbf{h} is the vector of
 133 unknown heads at time t ; and \mathbf{b} is a known vector of source/sink and storage terms.

134 3.1 Conjugate Gradient Algorithm


135 The conjugate gradient (CG) iterative method is an efficient method for solving large sys-
 136 tems of linear equations having a square, symmetric, positive-definite coefficient matrix.
 137 Pseudocode for the preconditioned conjugate gradient method is given in **algorithm 1**.

Algorithm 1 Preconditioned Conjugate Gradient Method

```

1: Compute  $\mathbf{r}_0 = \mathbf{b} - \mathbf{A}\mathbf{h}_0$  for some initial guess  $\mathbf{h}_0$ 
2: for  $i = 1 \rightarrow \text{maxinner}$  do
3:   solve  $\mathbf{M}\mathbf{z}_{i-1} = \mathbf{r}_{i-1}$  ▷ apply preconditioner
4:    $\rho_{i-1} = \mathbf{r}_{i-1}^T \mathbf{z}_{i-1}$ 
5:   if  $i = 1$  then
6:      $\mathbf{p}_i = \mathbf{z}_0$ 
7:   else
8:      $\beta_{i-1} = \rho_{i-1} / \rho_{i-2}$ 
9:      $\mathbf{p}_i = \mathbf{z}_{i-1} + \beta_{i-1} \mathbf{p}_{i-1}$ 
10:  end if
11:   $\mathbf{q}_i = \mathbf{A}\mathbf{p}_i$ 
12:   $\alpha_i = \rho_{i-1} / \mathbf{p}_i^T \mathbf{q}_i$  ▷ calculate step-size
13:   $\mathbf{h}_i = \mathbf{h}_{i-1} + \alpha_i \mathbf{p}_i$  ▷ update head
14:   $\mathbf{r}_i = \mathbf{r}_{i-1} - \alpha_i \mathbf{q}_i$  ▷ update residual
15:  check convergence; continue if necessary
16: end for

```

138 In **algorithm 1**, \mathbf{r} is the residual of **equation 4** [L^3/T], i is the linear iteration index
139 [unitless], *maxinner* is a user-specified maximum number of iterations to perform [unitless],
140 \mathbf{M} is the preconditioned form of \mathbf{A} [L^2/T], \mathbf{z} is the solution resulting from application of \mathbf{M}
141 to \mathbf{r} [L^3/T], \mathbf{p} is the orthogonal search direction [L], β is a scalar used to determine the next
142 search direction [unitless], \mathbf{q} is the residual change resulting from multiplication of \mathbf{A} and
143 the search direction \mathbf{p} [L^3/T], α is the step-size [unitless]. 

144 All linear iteration steps, except for application of the zero fill-in incomplete LU (ILU0)
145 and MILU0 preconditioners in line 3 of **algorithm 1**, include operations that are indepen-
146 dent for each active groundwater cell, indicating a potential for parallelization. In cases
147 where (1) no preconditioner ($\mathbf{z} = \mathbf{r}$) or (2) a Jacobi preconditioner ($\mathbf{M} = \mathbf{I}\mathbf{D}^{-1}$ – where \mathbf{I}
148 is the identity matrix and \mathbf{D}^{-1} is a matrix that only contains the inverse of the diagonal
149 elements of \mathbf{A}), is selected, application of the preconditioner is highly parallelizable. For
150 most practical problems, non-preconditioned and the Jacobi preconditioners will not signif-
151 icantly reduce the number of linear iterations (i) required to achieve convergence (Li and
152 Saad 2010). As a result better preconditioners, such as incomplete factorizations (ILU0 or
153 MILU0) or m-degree polynomial (GLSPOLY) approximations of \mathbf{A} , are generally needed.

Application of the incomplete factorization preconditioners requires a forward substitution ($\mathbf{L}\mathbf{y} = \mathbf{b}$ – where \mathbf{L} is the lower triangular portion of \mathbf{A}) followed by a backward substitution ($\mathbf{U}\mathbf{z} = \mathbf{y}$ – where \mathbf{U} is upper triangular portion of \mathbf{A}) and is difficult to efficiently parallelize (Barrett *et al.* 1994). Preconditioning using polynomial preconditioners, however, is simple, efficient, and very easy to parallelize (Liang *et al.* 2002; Saad 2003).

3.2 Numerical Implementation

The UPCG linear solver was coded to allow solution of the groundwater flow equation on either the CPU, GPGPU, or combination of CPU and GPGPU using double precision operations. Execution of the UPCG solver on the CPU, GPGPU, or combination of CPU and GPGPU is specified by the user in the UPCG input file.

Although the connectivity of the finite-difference discretization used in MODFLOW-2005 uses a fixed 5- and 7-point stencil in two- and three-dimensions, respectively, the PCG method in **algorithm 1** was coded using an unstructured compressed sparse row storage (CSR) format (Barrett *et al.* 1994) with the diagonal element in the first position in each row. This storage format is also supported by the CUBLAS libraries. The UPCG linear solver includes Jacobi, ILU0, MILU0, and GLSPOLY preconditioners (Saad 2003).

Available CUBLAS matrix-vector products and level-1 BLAS operations were used to parallelize the GPGPU capabilities of the UPCG linear solver. Because of the poor performance of forward and back substitution operations used by the ILU0 and MILU0 preconditioners observed during initial testing of the UPCG linear solver and observed by others (Li and Saad 2010), application of these preconditioners is performed on the CPU, resulting in a hybrid CPU/GPGPU solver. The results of the matrix-vector product, calculated by application of the ILU0 and MILU0 preconditioners on the CPU, is accessed as page locked memory by the GPGPU to achieve higher memory bandwidth.

For parallel CPU simulations, OpenMP directives were added to all matrix-vector products and level-1 BLAS operations. The OpenMP matrix-vector product was constructed

using a level-2 BLAS GEMV approach because of the slight increase in performance of this approach over dot-product and `axpy` approaches for calculating matrix vector products (Petersen and Arbenz 2004). OpenMP directives `NUM_THREADS` and `SCHEDULE STATIC` were used in all OpenMP routines to allow the user to control the number of OpenMP threads and reduce OpenMP scheduling/synchronization overhead, respectively. The number of OpenMP threads used for matrix-vector products and level-1 BLAS operations is specified separately in the UPGP input file. For sequential CPU simulations, separate matrix-vector product and level-1 BLAS operation routines that exclude OpenMP directives are used to eliminate all OpenMP overhead.

Infinity norms ($\|x\|_\infty$) are used to evaluate convergence of the UPGP linear solver with respect to the change in \mathbf{h} (HCLOSE) and \mathbf{r} (RCLOSE). In MODFLOW-2005, non-linearities are resolved using Picard iteration. Picard iteration is implemented such that the \mathbf{A} matrix is formulated using the latest estimate of \mathbf{h} , then \mathbf{h} and \mathbf{r} are updated using a linear solver. This process is continued until the maximum change in h and r in any model cell is less than HCLOSE and RCLOSE, respectively, on the first iteration ($i=1$) of the linear solver or the maximum number of outer (Picard) iterations is exceeded.

4 Test Cases

The performance of the GPGPU implementation of the UPGP linear solver was evaluated on a NVIDIA Tesla C2050 GPGPU with 3 GB of RAM and capable of 515 GFLOP/sec of double precision processing performance (NVIDIA compute capability 2.0). The GPGPU was mounted in a 16-pin PCI local bus on a single quad core Intel Xenon 3GHz CPU, capable of executing 4 threads per CPU, with 6 GB of RAM and running a 64-bit version of the Windows 7 Enterprise OS. For comparison with GPGPU results, parallel CPU simulations were executed on dual 4-core Intel Xenon 2.4 GHz CPUs, capable of executing 8 parallel threads, with 16 GB of RAM running a 64-bit version of the Windows Server 2008 R2

Standard OS (Service Pack 1).

4.1 Problem Description

GPGPU and parallel CPU results were evaluated using a number of horizontal and vertical model discretizations of a 1,000 m square heterogeneous unconfined aquifer. The top and bottom of the aquifer are flat and specified to have elevations of 10 and -30 m, respectively. Models with horizontal discretizations of 200×200 , 500×500 , $1,000 \times 1,000$, $2,000 \times 2,000$, and $4,000 \times 4,000$ cells per layer were evaluated. Single layer models were evaluated for all horizontal discretizations and multi-layer models were evaluated for select horizontal discretizations to determine the effect that a 7-point stencil, and the resulting additional non-zero entries for each cell, has on GPGPU and parallel CPU execution time. Two- and three-layer models were evaluated for all horizontal discretizations except $4,000 \times 4,000$ cells per layer. Ten-layer models were also evaluated for horizontal discretizations ranging from 200×200 to $1,000 \times 1,000$ cells per layer. Additional horizontal and vertical discretizations

were not evaluated because of GPGPU memory requirements for larger problems. For multi-layer models, a constant layer thickness was used for all but the first layer and was calculated by dividing 30 m by the number of layers. The thickness of layer 1 was calculated to be sum of the top elevation of the model (10 m) and constant thickness used for layers 2 and below.

The heterogeneous hydraulic conductivity distribution (base hydraulic conductivity data) was generated for the $1,000 \times 1,000$ cell model (base model) using sequential Gaussian simulation with log-transformed (base 10) parameters based on a mean value of 5 m/d and an exponential variogram having a nugget of $1.023 \log_{10}((m/d)^2)$, a variance of $1.23 \log_{10}((m/d)^2)$, and length parameter (a) of 250.5 m (effective range ≈ 750 m). The axis of anisotropy was rotated by 45 degrees to the model grid and a 3.5:1 ratio of horizontal anisotropy was used. For the realization selected and used to evaluate GPGPU and parallel CPU results, the minimum, average, and maximum hydraulic conductivity in the model domain were 0.43, 4.81, and 43.3 m/d, respectively. The distribution of hydraulic conductivity in the model domain

is shown in **figure 2**.

Hydraulic conductivity was not varied between layers in the multi-layer models evaluated. For models having less than $1,000 \times 1,000$ cells per layer, the base hydraulic conductivity data was averaged for each base model cell contained in a coarse model cell. For models having more than $1,000 \times 1,000$ cells per layer, the base hydraulic conductivity data was directly assigned to each higher-resolution cell contained within a coarse base model cell.

Head-dependent (general head) boundaries were specified **on the left and right sides** and no flow boundaries at the top and bottom. Head-dependent boundaries are assumed to be located a half model cell away from the edge of the model domain and conductance values were calculated using the average hydraulic conductivity of 4.81 m/d, vertical saturated cell areas, and half the model cell dimensions in the x-direction; head-dependent boundary head values were calculated using the hydraulic gradient over the model domain $[(10. - 0.0) / 1000. = 0.01 \text{ m/m}]$ and assumed heads at the left and right sides of the model domain of 10.0 and 0.0 m, respectively. For the one layer $1,000 \times 1,000$ cell model, conductance and head values of $214.1447 \text{ m}^2/\text{d}$, $160.6085 \text{ m}^2/\text{d}$, 10.025 m, and -0.025 m were specified for the left and right side of the model domain, respectively. The head-dependent boundaries cause an ambient groundwater flow from left to right. Four pumping wells, located in the four cells in the center of the model domain, pump at a constant rate with a total groundwater **withdrawal rate of $1,000 \text{ m}^3/\text{d}$** . For multi-layer simulations, the pumping wells withdraw water from the lower model layer. There is no recharge or evapotranspiration; all groundwater flow to the wells is provided by the head-dependent boundaries. Boundary conditions are shown graphically in **figure 2**.

For all model discretizations, the UPGP solver with the Jacobi, MILU0, and GLSPOLY preconditioners were evaluated on the CPU, using both the sequential and parallel (OpenMP) options, and GPGPU. To maximize the amount of time spent in the linear solvers the maximum number of outer (Picard) and inner (linear) iterations were set to 50 and 1,000, respectively. A HCLOSE value of 0.001 m was used for all simulations. In order to use

a consistent RCLOSE value in all simulations, RCLOSE was calculated as the product of HCLOSE and the cell area, effectively scaling RCLOSE with the discretization. A degree 10 polynomial was constructed for simulations using the GLSPOLY preconditioner and because the coefficient matrix \mathbf{A} was scaled using diagonal scaling we assumed the minimum and maximum eigenvalues for the scaled matrix were between 0.0 and 2.0 (Scandrett 1989).

All of the models simulated steady-state conditions. As a result, specific storage and specific yield values were not specified for the aquifer. Initial heads were specified to be 0.0 m throughout the model domain, sufficiently different from final steady-state heads. Simulated heads for the one layer $1,000 \times 1,000$ cell model are shown in **figure 2** and show the effect that groundwater withdrawals and heterogeneous hydraulic conductivity have on water levels.

4.2 Results

The CPU/GPGPU speedup of UPGP solver simulations for the Jacobi, MILU0, and GLSPOLY preconditioners are summarized on **table 1**. Significant CPU/GPGPU speedups, ranging from 1.5 to 35, are observed with Jacobi and GLSPOLY preconditioners. Only marginal speedups, ranging from less than 1.0 to almost 2.0, are observed with the MILU0 preconditioner. The maximum number of outer (Picard) and inner (linear) iterations were exceeded for the $4,000 \times 4,000 \times 1$ cell problem simulated using the Jacobi preconditioner, highlighting that the Jacobi preconditioner is not a robust preconditioner for some problems. The smallest CPU/GPGPU speedup corresponds with the smallest model discretizations ($200 \times 200 \times 1$) because of the overhead associated with GPGPU-CPU memory copy operations relative to the solution time.

The CPU/GPGPU speedup of MILU0 and GLSPOLY simulations are shown graphically on **figure 3**. The CPU/GPGPU speedup of MILU0 simulations are approximately 2 for simulations with more than 3 million active model cells. The reduction in the CPU/GPGPU speedup of MILU0 simulations between 1 and 3 million active model cells is a result of

the higher percentage of runtime spent transferring preconditioner data (\mathbf{r} and \mathbf{z} in **algorithm 1**) between the GPGPU and CPU during each linear iteration for application of the MILU0 preconditioner. The CPU/GPGPU speedup of GLSPOLY preconditioner increases to approximately 30 for simulations with more than 5 million active cells; the increase in the CPU/GPGPU speedup with increased problem size for the Jacobi preconditioner is comparable to the GLSPOLY preconditioner (**fig. 3b**).

The amount of time applying the MILU0 and GLSPOLY preconditioners, performing matrix-vector operations, level-1 BLAS operations, and transferring data between the CPU and GPU for a select number of simulations are shown graphically in **figure 4**. Matrix-vector operations account for approximately half of the time spent in the UPGP solver for sequential CPU simulations using either the MILU0 or GLSPOLY preconditioners. On the GPGPU, matrix-vector operations for the MILU0 and GLSPOLY preconditioner are significantly reduced, largely as a result of the ability of the GPGPU to apply massively parallel resources to perform BLAS operations. Analysis of operation execution times indicates that the MILU0 speedup is notably less than the GLSPOLY speedup because the time spent applying the MILU0 is essentially the same whether the GPGPU is used or not. This is because the application of the MILU0 preconditioner is always performed on the CPU. For the MILU0 preconditioner, all of the speedup occurs during the conjugate gradient routine in the level-1 BLAS operations [matrix-vector multiplication and vector operations (dot products)], which are carried out on the GPGPU after application of the preconditioner. Furthermore, because of additional memory transfers between the CPU and GPGPU when using the MILU0 preconditioner, the total time spent applying the preconditioner actually increases. For the GLSPOLY preconditioner, solution on the GPGPU reduces the time spent completing linear solver operations by more than 95% and the time spent transferring data between the CPU and GPGPU is negligible.

5 Discussion

Although notable, CPU/GPGPU speedups summarized on **table 1** are not a fair measure of the speedups that would result from use of the GPGPU to solve a MODFLOW-2005 problem. A more realistic comparison is the speedup of GPGPU simulations relative to sequential CPU simulations using the standard MODFLOW-2005 PCG solver with the modified incomplete Cholesky (MIC) preconditioner (Hill 1990). PCG(MIC)/GPGPU speedups for the Jacobi, MILU0, and GLSPOLY preconditioners are summarized on **table 2**. PCG(MIC)/GPGPU speedups for the MILU0 and GLSPOLY preconditioners are shown graphically on **figure 5**.

PCG(MIC)/GPGPU speedups for the MILU0 preconditioner are comparable to CPU/GPGPU speedups; the similarity of PCG(MIC)/GPGPU and CPU/GPGPU speedups for the MILU0 preconditioner are a result of the similarity in time required to perform forward and backward substitutions during application of the preconditioner with the PCG and UPGP solvers, as the MIC is similar to MILU0. PCG(MIC)/GPGPU speedups for the Jacobi and GLSPOLY preconditioners range from less than 1.0 to more than 8.0, which are notably less than CPU/GPGPU speedups. This is attributable to the efficiency of the standard MODFLOW-2005 PCG solver with the MIC preconditioner. The relative increase in PCG(MIC)/GPGPU speedups for the Jacobi and GLSPOLY preconditioners is a result of the reduced efficiency of the PCG solver with the MIC preconditioner as the number of model layers increase (**fig. 5b**).

Parallel CPU simulations were run to compare speedups possible with GPGPUs to speedups that could be realized on multi-core CPUs using the OpenMP programming paradigm. PCG(MIC)/UPGP-OpenMP speedups for the Jacobi, MILU0, and GLSPOLY preconditioners using 4, 7, and 14 OpenMP threads are summarized on **table 3**. PCG(MIC)/UPGP-OpenMP speedups for the Jacobi and GLSPOLY preconditioners are notably less than PCG(MIC)/GPGPU speedups and is expected because the number of threads on the GPGPU is orders of magnitude greater than for parallel CPU simulations. For the MILU0 preconditioner, PCG(MIC)/UPGP-OpenMP speedups (average speedup = 1.539) are comparable

to PCG(MIC)/GPGPU speedups (average speedup = 1.447) and there is little benefit to using the GPGPU with the MILU0 preconditioner. PCG(MIC)/UPCG-OpenMP speedups for the MILU0 preconditioner with 7 OpenMP threads are shown graphically on **figure 6**. It should be noted that the Intel Fortran compiler (version 12.1.2.278), with the highest-level optimization option (/O3 compiler switch), was used to compile MODFLOW-2005 with the UPCG solver. The relatively small PCG(MIC)/UPCG-OpenMP speedups observed with the UPCG solver compiled using the Intel Fortran compiler is consistent with the observations of Dong and Li (2009). Larger PCG(MIC)/UPCG-OpenMP speedups would likely be observed with other compilers but the value of using a less optimized compiler is questionable.

6 Conclusions

The PCG algorithm was implemented for use on GPGPUs using the NVIDIA CUDA-enabled implementation of the BLAS library (CUBLAS). The availability of this API makes it relatively simple to develop parallel GPGPU code. The fork-join parallel model, which uses multiple threads in parallel sections, is used to parallelize code on GPGPUs. Unlike fork-join approaches used on CPUs (OpenMP), which have thread limits determined by the number of CPUs and cores per CPU, GPGPUs have the capability of using thousands to tens of thousands of threads concurrently in parallel sections. The number of concurrent threads that can be used on GPGPUs makes them ideal for parallelization of simultaneous solutions of systems of linear equations like those solved by MODFLOW-2005. The key to successful parallelization on GPGPUs is to (1) use numerical approaches that are highly parallelizable and (2) reduce the need for memory transfers between the CPU and GPGPU as much as possible.

The UPCG solver was developed to solve the PCG algorithm on GPGPUs and includes Jacobi, ILU0, MILU0, and GLSPOLY preconditioners. An unstructured CSR format was used in the UPCG solver to facilitate use of the CUBLAS library. The UPCG solver also

includes options for running sequential and parallel simulations on the CPU. Parallel CPU simulations are executed using the OpenMP programming paradigm.

A number of test problems were evaluated and indicate that it is possible to realize notable speedup through use of GPGPUs with MODFLOW-2005. CPU/GPGPU speedups ranging from 1.5 to 35 were observed for the Jacobi, MILU0, and GLSPOLY preconditioners when GPGPU runtimes were compared to CPU runtimes for the same preconditioner. Exceptions to notable CPU/GPGPU speedups were observed for the smallest problems where CPU-GPGPU memory transfer times were significant compared to solution time. CPU/GPGPU speedups with the MILU0 preconditioner was the smallest of all of the preconditioners (ranging from 0.99 to 1.936) because of the additional CPU-GPGPU memory transfers required during each inner (linear) iteration to apply the preconditioner.

Speedups calculated using GPGPU runtimes and CPU runtimes for the same solver are not fair measures of the benefits of using GPGPUs to solve the groundwater flow equation because of the efficiency of the standard MODFLOW-2005 PCG solver with the MIC preconditioner. PCG(MIC)/GPGPU speedups are smaller than speedups calculated using sequential CPU runtimes for the UPCG solver. Use of GPGPUs with the UPCG solver and the GLSPOLY preconditioner can result in PCG(MIC)/GPGPU speedups as large as 8 for large multi-layer models when compared to the standard MODFLOW-2005 PCG solver with the MIC preconditioner. For certain problems, the Jacobi preconditioner is adequate and has similar PCG(MIC)/GPGPU speedups.

When compared to parallel CPU solutions, speedups for GPGPU solutions using the UPCG solver with the Jacobi and GLSPOLY preconditioners are significantly better. For problems where the MILU0 preconditioner is needed there is no real benefit to using the GPGPU because of the required additional CPU-GPGPU memory transfers – for these cases parallel CPU solution would be adequate.

In summary, use of GPGPU results in significant speedups when the Jacobi and GLSPOLY preconditioners are capable of reducing the total number of inner (linear) iterations required

to achieve specified convergence criteria. The benefits of using the GPGPU increases as the problem size increases. In this study, problem size was limited by the RAM available on the GPGPU. Larger problems could be solved with GPGPUs having more RAM or extending the UPGC solver to allow use of multiple GPGPUs. The UPGC solver is designed to solve square, symmetric, positive-definite coefficient matrices. Modification of the UPGC solver to allow solution of nonsymmetric matrices, using the BiConjugate Gradient Stabilized method or other similar method (Barrett *et al.* 1994), would permit use of the UPGC solver with MODFLOW-NWT (Niswonger *et al.* 2011), MT3DMS (Zheng and Wang 1999), or SEAWAT (Langevin *et al.* 2007).

7 Acknowledgements

The authors are grateful to Randall T. Hanson and Yong Liu for the valuable suggestions and helpful comments on previous versions of this article.

8 Supporting Information

Additional Supporting Information may be found in the online version of this article:

The supporting information includes source code for the UPGC solver and the MODFLOW-2005 main routine.

Please note: Wiley-Blackwell is not responsible for the content or functionality of any supporting information supplied by the authors. Any queries (other than missing material) should be directed to the corresponding author for the article.

9 References

- Barrett R., Berry, M., Chan, T.F., Demmel, J., Donato, J., Dongarra, J., Eijkhout, V.,
 Pozo, R., Romine, C., Van Der Vorst, H. 1994. Templates for the Solution of Linear
 Systems, Building Blocks for Iterative Methods. Philadelphia, Pennsylvania, SIAM,
 124 p.
- Dong, Yanhui, Li, Guomin 2009. A Parallel PCG Solver for MODFLOW. Ground Water
 47: 845–850. doi: 10.1111/j.1745-6584.2009.00598.x
- Harbaugh, A.W. 2005. MODFLOW-2005, the U.S. Geological Survey modular ground-
 water model – the Ground-Water Flow Process. U.S. Geological Survey Techniques
 and Methods 6-A16, variously p.
- Hill, M.C. 1990. Preconditioned conjugate-gradient 2 (PCG2), a computer program for
 solving ground-water flow equations. U.S. Geological Survey Water-Resources Investi-
 gations Report 90-4048, 43 p.
- Ji, Xiaohui, Cheng, Tangpei, Wang, Qun 2012. CUDA-based solver for large-scale ground-
 water flow simulation. Engineering with Computers 28, no. 1: 13-19. doi: 10.1007/s00366-
 011-0213-2
- Langevin, C.D., Thorne, D.T., Jr., Dausman, A.M., Sukop, M.C., and Guo, Weixing 2007.
 SEAWAT Version 4: A Computer Program for Simulation of Multi-Species Solute and
 Heat Transport. U.S. Geological Survey Techniques and Methods Book 6, Chapter
 A22, 39 p.
- Li, Ruipeng, Saad, Yousef 2010. GPU-accelerated preconditioned iterative linear solvers.
 Minnesota Supercomputer Institute Report umsi-2010-112, University of Minnesota,
 Minneapolis, MN, 24 p.

- Liang, Y., Weston, J., and Szularz, M. 2002. Generalized least-squares polynomial preconditioner for symmetric indefinite linear equations, *Parallel Computing* 28: 323-341.
- McDonald, M.G., Harbaugh, A.W. 1988. A modular three-dimensional finite-difference ground-water flow model. U.S. Geological Survey Techniques of Water-Resources Investigations, book 6, chap. A1, 586 p.
- Mehl, S.E., Hill, M.C. 2001. MODFLOW-2000, the U.S. Geological Survey modular ground-water model – User guide to the LINK-AMG (LMG) Package for solving matrix equations using an algebraic multigrid solver. U.S. Geological Survey Open-File Report 01-177, 33 p.
- Naff, R.L. 2008. Technique and Application of a Parallel Solver to MODFLOW, *Proceedings of MODFLOW and More 2008: Ground Water and Public Policy*, v. 5, May 19-21, 2008, Colorado School of Mines, Golden, Colorado.
- Naff, R.L., Banta, E.R. 2008. The U.S. Geological Survey modular ground-water model—PCGN: A preconditioned conjugate gradient solver with improved nonlinear control. U.S. Geological Survey Open-File Report 2008–1331, 35 p.
- Niswonger, R.G., Panday, Sorab, and Ibaraki, Motomu 2011. MODFLOW-NWT, A Newton formulation for MODFLOW-2005. U.S. Geological Survey Techniques and Methods 6-A37, 44 p.
- NVIDIA 2011. NVIDIA CUDA™: NVIDIA CUDA C Programming Guide Version 4.1.
- OpenMP Architecture Review Board. 2005. OpenMP Application Program Interface, Version 2.5, accessed March 17, 2011 at <http://www.openmp.org/mp-documents/spec25.pdf>.
- Petersen, W.P. and Arbenz, Peter 2004. *Introduction to Parallel Computing: A Practical Guide with Examples in C*. New York, Oxford University Press, 288 p.

- 456 Saad, Yousef 2003. Iterative methods for sparse linear systems – 2nd edition, Philadelphia,
457 Pennsylvania, SIAM, 528 p.
- 458 Scandrett, C. 1989. Comparison of several iterative techniques in the solution of symmetric
459 banded equations on a two-pipe Cyber 250, Applied Mathematics and Computation
460 34: 95-112.
- 461 Soni, Kunal, Chandar, D.D.J., and Sitaraman, J. 2012. Development of an overset grid
462 computational fluid dynamics solver on graphical processing units. Computers & Flu-
463 ids 58: 1-4. doi: 10.1016/j.compfluid.2011.11.014.
- 464 Wilson, J.D., Naff, R.L. 2004. The U.S. Geological Survey modular ground-water model –
465 GMG linear equation solver package documentation. U.S. Geological Survey Open-File
466 Report 2004-1261, 47 p.
- 467 Zheng, Chunmiao and Wang, P.P. 1999 MT3DMS. A modular three-dimensional multi-
468 species transport model for simulation of advection, dispersion and chemical reactions
469 of contaminants in groundwater systems; documentation and user's guide. U.S. Army
470 Engineer Research and Development Center Contract Report SERDP-99-1, Vicksburg,
471 MS, 202 p.

Table 1: Speedup of UPCG solver GPGPU simulations relative to CPU simulations for the same UPCG preconditioner. Speedup values in **bold** type indicate simulations that did not converge within 50 outer iterations with up to 1,000 inner iterations.

Columns \times Rows \times Layers	Speedup		
	Jacobi	MILU0	GLSPOLY
$200 \times 200 \times 1$	1.513	0.990	2.972
$500 \times 500 \times 1$	9.536	1.699	16.315
$1000 \times 1000 \times 1$	18.871	1.404	26.786
$2000 \times 2000 \times 1$	24.453	1.938	33.615
$4000 \times 4000 \times 1$	26.918	1.936	35.154
$200 \times 200 \times 2$	3.234	1.517	6.146
$500 \times 500 \times 2$	13.420	1.742	18.665
$1000 \times 1000 \times 2$	19.877	1.322	25.420
$2000 \times 2000 \times 2$	23.871	1.861	29.516
$200 \times 200 \times 3$	5.234	1.517	9.231
$500 \times 500 \times 3$	15.848	1.769	20.843
$1000 \times 1000 \times 3$	22.444	1.265	27.658
$2000 \times 2000 \times 3$	25.171	1.870	30.760
$200 \times 200 \times 10$	11.552	1.715	17.564
$500 \times 500 \times 10$	20.806	1.833	25.826
$1000 \times 1000 \times 10$	24.981	1.889	29.888

Table 2: Speedup of UPCG solver GPGPU simulations relative to simulations performed using the PCG solver with the modified incomplete Cholesky (MIC) preconditioner. CPU times for the PCG solver with the MIC preconditioner are in seconds. Speedup values in **bold** type indicate simulations that did not converge within 50 outer iterations with up to 1,000 inner iterations.

Columns \times Rows \times Layers	CPU Time	Speedup		
	MIC	Jacobi	MILU0	GLSPOLY
200 \times 200 \times 1	0.6826875	0.487	1.101	0.841
500 \times 500 \times 1	10.27194	1.675	1.543	2.294
1000 \times 1000 \times 1	75.64222	3.241	1.239	2.957
2000 \times 2000 \times 1	536.2608	1.891	1.438	2.754
4000 \times 4000 \times 1	3891.683	1.357	1.696	2.839
200 \times 200 \times 2	1.683125	1.081	1.330	1.633
500 \times 500 \times 2	24.66269	3.156	1.534	3.497
1000 \times 1000 \times 2	175.7042	4.111	1.156	3.532
2000 \times 2000 \times 2	1058.137	1.967	1.590	2.837
200 \times 200 \times 3	3.040594	1.875	1.487	2.349
500 \times 500 \times 3	38.59184	3.908	1.530	4.075
1000 \times 1000 \times 3	315.6253	5.105	1.130	4.428
2000 \times 2000 \times 3	1962.365	2.560	1.634	3.562
200 \times 200 \times 10	18.50184	5.344	1.524	4.921
500 \times 500 \times 10	218.9169	8.683	1.598	7.643
1000 \times 1000 \times 10	1621.158	9.299	1.621	8.633

Table 3: Speedup of UPCG solver OpenMP simulations, with 4, 7, and 14 threads, relative to simulations performed using the PCG solver with the modified incomplete Cholesky (MIC) preconditioner. Speedup values in **bold** type indicate simulations that did not converge within 50 outer iterations with up to 1,000 inner iterations.

Columns \times Rows \times Layers	Speedup – 4 / 7 /14 Threads		
	Jacobi	MILU0	GLSPOLY
200 \times 200 \times 1	1.166/0.614/0.614	2.446/1.511/1.423	1.185/0.779/1.241
500 \times 500 \times 1	0.447/0.332/0.357	1.637/1.197/1.164	0.571/0.364/0.535
1000 \times 1000 \times 1	0.409/0.419/0.359	1.488/1.463/1.194	0.313/0.323/0.309
2000 \times 2000 \times 1	0.185/0.182/0.151	1.691/1.622/1.286	0.225/0.272/0.198
4000 \times 4000 \times 1	0.111/0.099/0.108	1.619/1.523/1.670	0.222/0.241/0.245
200 \times 200 \times 2	1.113/1.090/1.004	1.954/2.258/1.567	1.171/1.147/1.626
500 \times 500 \times 2	0.536/0.567/0.460	1.500/1.521/1.151	0.525/0.625/0.556
1000 \times 1000 \times 2	0.527/0.472/0.411	1.543/1.474/1.227	0.396/0.295/0.335
2000 \times 2000 \times 2	0.211/0.198/0.168	1.401/1.693/1.418	0.237/0.313/0.244
200 \times 200 \times 3	0.872/1.067/0.897	1.433/1.901/1.300	0.827/1.080/1.346
500 \times 500 \times 3	0.577/0.482/0.527	1.444/1.417/1.328	0.548/0.579/0.598
1000 \times 1000 \times 3	0.593/0.562/0.479	1.519/1.590/1.377	0.389/0.340/0.385
2000 \times 2000 \times 3	0.245/0.246/0.215	1.746/1.478/1.544	0.313/0.286/0.295
200 \times 200 \times 10	1.066/1.059/1.215	1.503/1.574/1.426	0.858/0.899/1.186
500 \times 500 \times 10	0.970/1.014/1.019	1.563/1.604/1.400	0.780/0.799/0.913
1000 \times 1000 \times 10	0.835/0.968/0.928	1.720/1.734/1.651	0.815/0.798/0.914

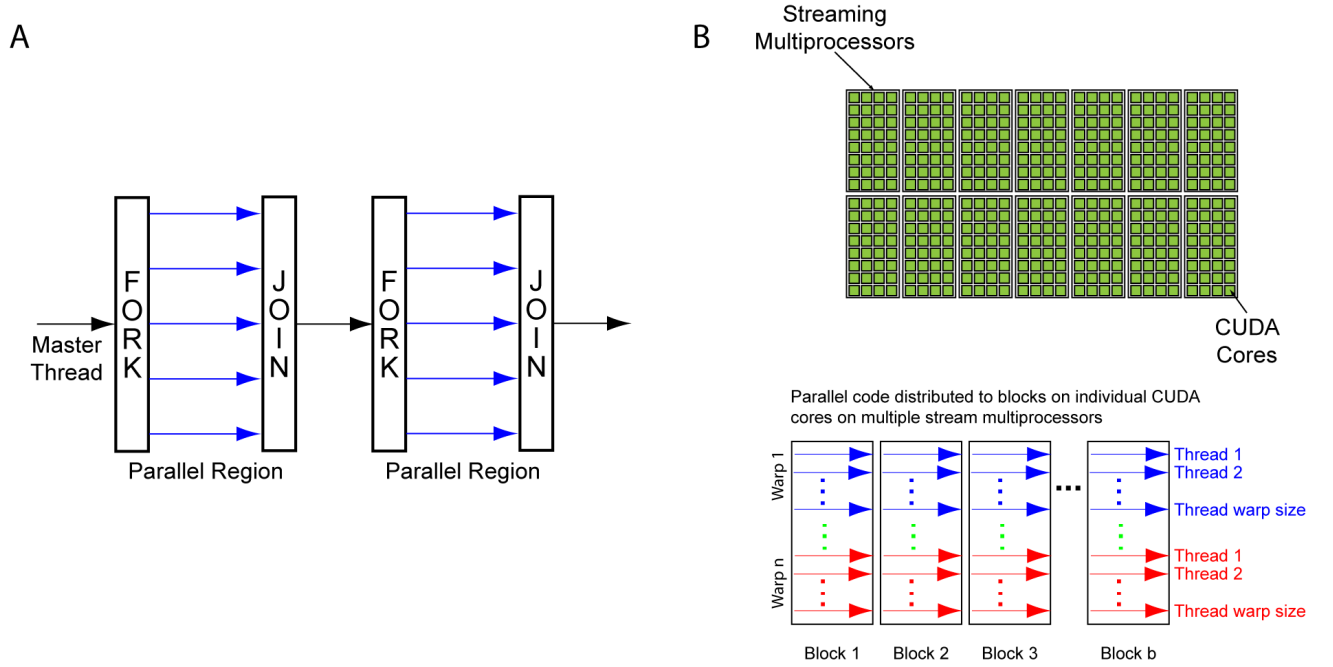


Figure 1: (A) Fork-join parallel model. The master thread executes sequential operations and initiates thread forks used in parallel regions. (B) GPGPU architecture showing the relation of streaming multiprocessors and CUDA cores. The fork-join parallel model is applied to blocks of code being executed by thread processing units on multiple streaming multiprocessors.




Figure 2: Model domain showing the distribution of hydraulic conductivity and boundary conditions applied to a hypothetical unconfined aquifer. A total of $1,000 \text{ m}^3/\text{day}$ are withdrawn from 4 pumping wells. Steady-state groundwater head contours (1 m contour interval) are also shown. Note the effect that groundwater withdrawals and hydraulic conductivity are having on simulated groundwater heads.




Figure 3: Speedup of UPGC solver GPGPU simulations with the (A) MILU0 and (B) GLSPOLY preconditioners relative to sequential UPGC solver simulations executed on the CPU. Note the different scales used for MILU0 and GLSPOLY speedup.




Figure 4: Time, in seconds, spent applying the preconditioner, performing BLAS level 1 operations, and transferring data between the CPU and GPGPU for (A) $2,000 \times 2,000 \times 3$ problem using the MILU0 preconditioner, (B) $4,000 \times 4,000 \times 1$ problem using the MILU0 preconditioner, (C) $2,000 \times 2,000 \times 3$ problem using the GLSPOLY preconditioner, and (D) $4,000 \times 4,000 \times 1$ problem using the GLSPOLY preconditioner. Note operation execution times have been normalized using the maximum preconditioner application time for each problem to facilitate comparison of different model runs.




Figure 5: Speedup of UPGC solver GPGPU simulations with the (A) MILU0 and (B) GLSPOLY preconditioners relative to simulations performed using the PCG solver with the modified incomplete Cholesky (MIC) preconditioner.

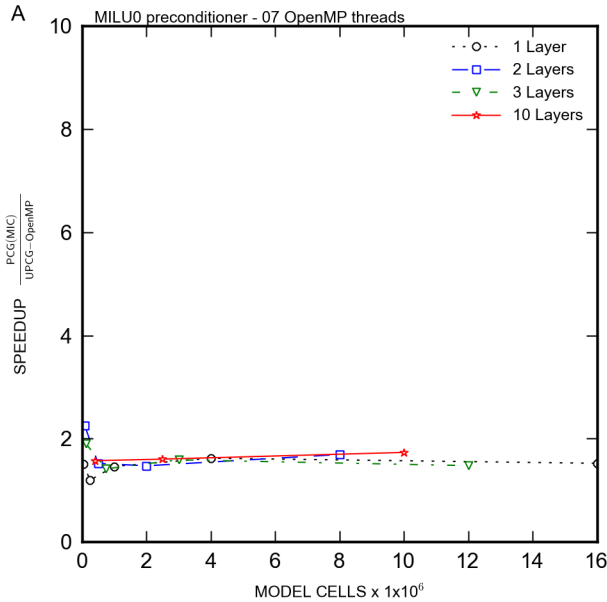


Figure 6: Speedup of parallel CPU simulations using the UPGC solver with the MILU0 preconditioner and 7 threads relative to simulations performed using the PCG solver with the modified incomplete Cholesky (MIC) preconditioner.